



# KUBERNETES ARCHITECTURE

kube-apiserver



Master

Manage, Plan, Schedule, Monitor  
Nodes



Worker Nodes

Host Application as Containers

kubelet

Controller-  
Manager


D  
CLUSTER

e-scheduler

Kube-proxy

Container Runtime Engine

Run containers



# Kubernetes Architecture



## Master

Manage, Plan, Schedule, Monitor  
Nodes

ETCD  
CLUSTER

kube-  
apiserver

Kube  
Controller  
Manager

kube-scheduler



## Worker Nodes

Host Application as Containers

kubelet

Kube-proxy

Container Runtime Engine

Run containers



kubelet

Kube-proxy

Container Runtime Engine

Run containers



# ETCD

## FOR BEGINNERS





**ETCD is a distributed  
reliable key-value store  
that is Simple, Secure &  
Fast**

# Operate ETCD

## 3. Run ETCD Service

```
./etcd
```

```
./etcdctl set key1 value1
```

```
./etcdctl get key1
```

```
value1
```

```
./etcdctl
```

NAME:

etcdctl - A simple command line client for etcd.

COMMANDS:

backup	backup an etcd directory
cluster-health	check the health of the etcd cluster
mk	make a new key with a given value
mkdir	make a new directory
rm	remove a key or a directory
rmdir	removes the key if it is an empty directory or a key-value pair
get	retrieve the value of a key


*Install from binary dir at top*



# ETCD

## In Kubernetes

# ETCD in HA Environment



etcd.service

```
ExecStart=/usr/local/bin/etcd \\  
  --name ${ETCD_NAME} \\  
  --cert-file=/etc/etcd/kubernetes.pem \\  
  --key-file=/etc/etcd/kubernetes-key.pem \\  
  --peer-cert-file=/etc/etcd/kubernetes.pem \\  
  --peer-key-file=/etc/etcd/kubernetes-key.pem \\  
  --trusted-ca-file=/etc/etcd/ca.pem \\  
  --peer-trusted-ca-file=/etc/etcd/ca.pem \\  
  --peer-client-cert-auth \\  
  --client-cert-auth \\  
  --initial-advertise-peer-urls https://${INTERNAL_IP}:2380 \\  
  --listen-peer-urls https://${INTERNAL_IP}:2380 \\  
  --listen-client-urls https://${INTERNAL_IP}:2379,https://127.0.0.1:2379 \\  
  --advertise-client-urls https://${INTERNAL_IP}:2379 \\  
  --initial-cluster-token etcd-cluster-0 \\  
  --initial-cluster controller-0=https://${CONTROLLER0_IP}:2380,controller-1=https://${CONTROLLER1_IP}:2380 \\  
  --initial-cluster-state new \\  
  --data-dir=/var/lib/etcd
```

*Important*



# kube-api server



# | Kube-api Server

1. Authenticate User

2. Validate Request

3. Retrieve data

4. Update ETCD

5. Scheduler

6. Kubelet

# Installing kube-api server

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kube-apiserver
```

## kube-apiserver.service

```
ExecStart=/usr/local/bin/kube-apiserver \\  
  --advertise-address=${INTERNAL_IP} \\  
  --allow-privileged=true \\  
  --apiserver-count=3 \\  
  --authorization-mode=Node,RBAC \\  
  --bind-address=0.0.0.0 \\  
  --enable-admission-  
plugins=Initializers,NamespaceLifecycle,NodeRestriction,LimitRanger,ServiceAccount,DefaultStorageClass,ResourceQuota \\  
  --enable-swagger-ui=true \\  
  --etcd-servers=https://127.0.0.1:2379 \\  
  --event-ttl=1h \\  
  --experimental-encryption-provider-config=/var/lib/kubernetes/encryption-config.yaml \\  
  --runtime-config=api/all \\  
  --service-account-key-file=/var/lib/kubernetes/service-account.pem \\  
  --service-cluster-ip-range=10.32.0.0/24 \\  
  --service-node-port-range=30000-32767 \\  
  --v=2
```

*Bind*  
*Link old server*  
*Used to verify service account tokens*

# Kube Controller Manager



# Controller

kube Api Server

Node 1

Node 2

Watch Status

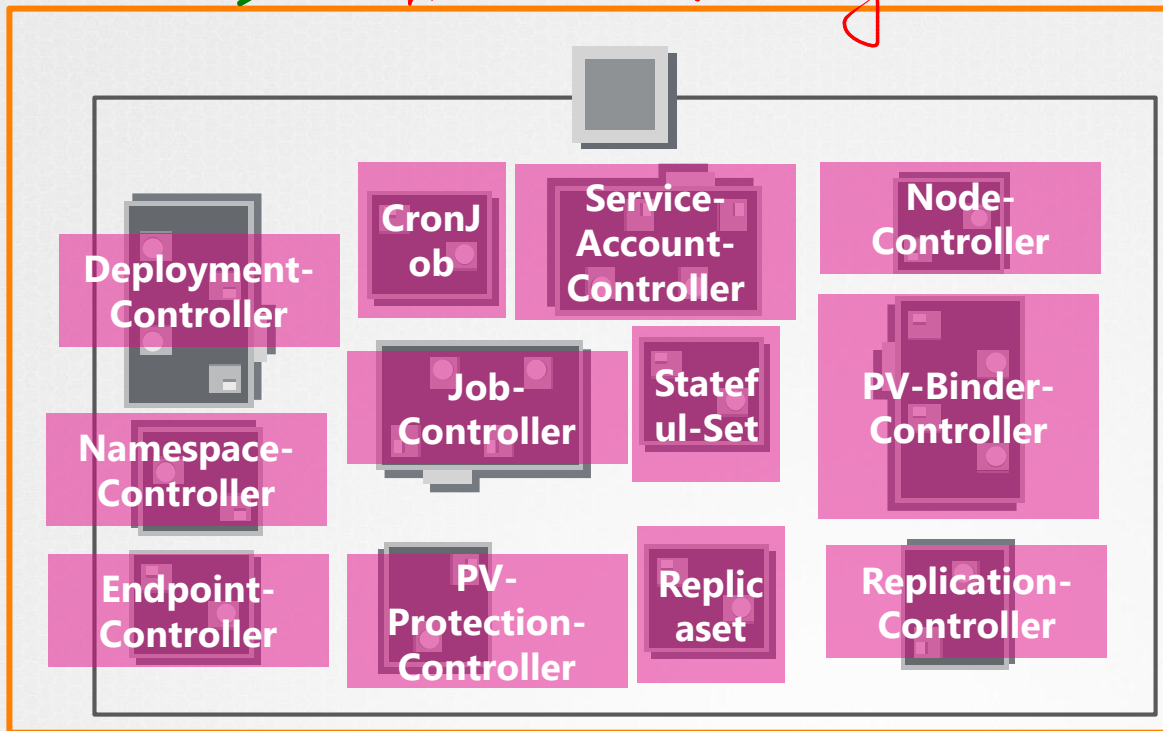
Remediate Situation

Node Monitor Period = 5s

Node Monitor Grace Period = 40s

POD Eviction Timeout = 5m

Kube Controller Manager



manage desired state of each object

After each 5 sec will check status

Controller waits for 40s for a response before marking it unhealthy

Time to wait for graceful termination

KODECLOUD

# Installing kube-controller-manager

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kube-controller-manager
```

## kube-controller-manager.service

```
ExecStart=/usr/local/bin/kube-controller-manager \\  
  --address=0.0.0.0 \\  
  --cluster-cidr=10.200.0.0/16 \\  
  --cluster-name=kubernetes \\  
  --cluster-signing-cert-file=/var/lib/kubernetes/ca.pem \\  
  --cluster-signing-key-file=/var/lib/kubernetes/ca-key.pem \\  
  --kubeconfig=/var/lib/kubernetes/kube-controller-manager.kubeconfig \\  
  --leader-elect=true \\  
  --root-ca-file=/var/lib/kubernetes/ca.pem \\  
  --service-account-private-key-file=/var/lib/kubernetes/service-account-key.pem \\  
  --service-cluster-ip-range=10.32.0.0/24 \\  
  --use-service-account-credentials=true \\  
  --v=2
```

```
N --node-monitor-period=5s
```

```
--controllers stringSlice      Default: [*]
```

A list of controllers to enable. '\*' enables all on-by-default controllers, 'foo' enables the controller named 'foo', '-foo' disables the controller named 'foo'.

All controllers: attachdetach, bootstrapsigner, clusterrole-aggregation, cronjob, csrapproving, csrcleaner, csrsigning, daemonset, deployment, disruption, endpoint, garbagecollector, horizontalpodautoscaling, job, namespace, nodeipam, nodelifecycle, persistentvolume-binder, persistentvolume-expander, podgc, pv-protection, pvc-protection, replicaset, replicationcontroller,

```
d=40s
```

```
s
```

# Installing kube-controller-manager

```
--controllers stringSlice      Default: [*]
```

A list of controllers to enable. '\*' enables all on-by-default controllers, 'foo' enables the controller named 'foo', '-foo' disables the controller named 'foo'.

All controllers: attachdetach, bootstrapsigner, clusterrole-aggregation, cronjob, csrapproving, csrcleaner, csrsigning, daemonset, deployment, disruption, endpoint, garbagecollector, horizontalpodautoscaling, job, namespace, nodeipam, nodelifecycle, persistentvolume-binder, persistentvolume-expander, podgc, pv-protection, pvc-protection, replicaset, replicationcontroller, resourcequota, root-ca-cert-publisher, route, service, serviceaccount, serviceaccount-token, statefulset, tokencleaner, ttl, ttl-after-finished


Disabled-by-default controllers: bootstrapsigner, tokencleaner

# Kube Scheduler







# Kube-Scheduler



1. Filter Nodes

2. Rank Nodes



# | More Later...

- Resource Requirements and Limits ✓
- Taints and Tolerations ✓
- Node Selectors/Affinity ✓

# Installing kube-scheduler

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kube-scheduler
```

kube-scheduler.service

```
ExecStart=/usr/local/bin/kube-scheduler \\  
--config=/etc/kubernetes/config/kube-scheduler.yaml \\  
--v=2
```

# View kube-scheduler options - kubeadm



```
cat /etc/kubernetes/manifests/kube-scheduler.yaml
```

```
spec:
  containers:
  - command:
    - kube-scheduler
    - --address=127.0.0.1
    - --kubeconfig=/etc/kubernetes/scheduler.conf
    - --leader-elect=true
```

# Kubelet

An abstract geometric pattern consisting of several interconnected lines and dots, forming a network-like structure, located in the bottom-left corner of the slide.

# Kubernetes Architecture



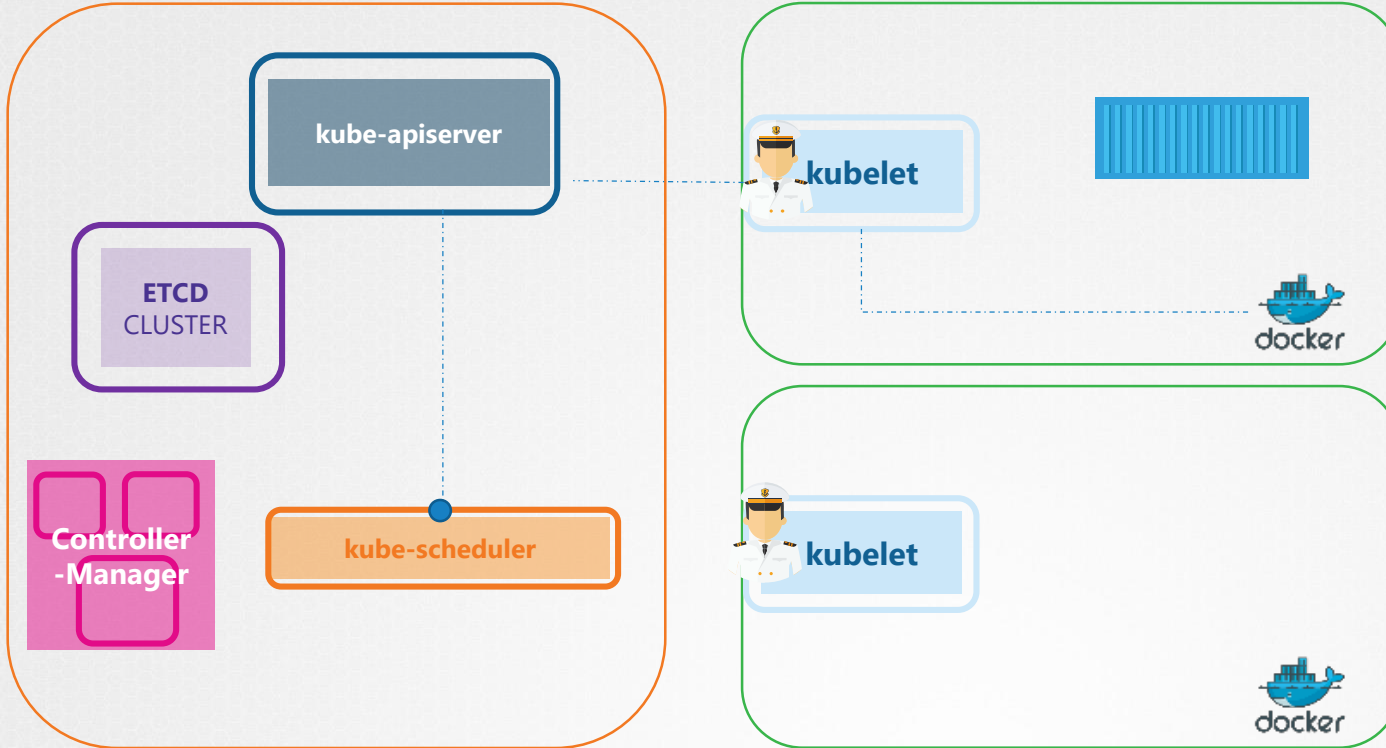
## Master

Manage, Plan, Schedule, Monitor  
Nodes



## Worker Nodes

Host Application as Containers



Register Node

Create PODs

Monitor Node & PODs

↓  
it also updates  
status of pods  
via REST api of  
kube-apiserver

# Installing kubelet

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kubelet
```

kubelet.service

```
ExecStart=/usr/local/bin/kubelet \\  
--config=/var/lib/kubelet/kubelet-config.yaml \\  
--container-runtime=remote \\  
--container-runtime-endpoint=unix:///var/run/containerd/containerd.sock \\  
--image-pull-progress-deadline=2m \\  
--kubeconfig=/var/lib/kubelet/kubeconfig \\  
--network-plugin=cni \\  
--register-node=true \\  
--v=2
```



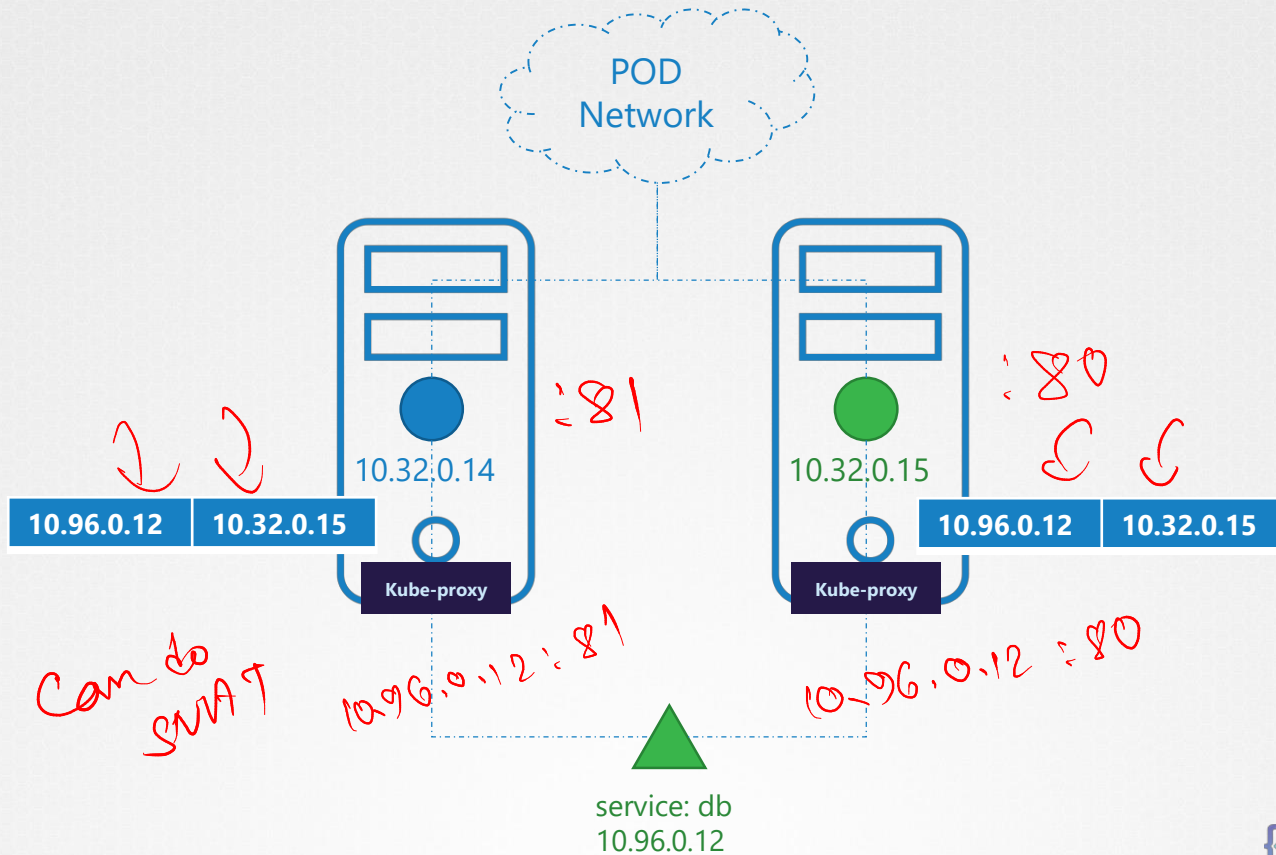
Kubeadm does not deploy  
Kubelets

# Kube-proxy

An abstract network diagram in the bottom-left corner, consisting of several orange dots (nodes) connected by thin orange lines, forming a web-like structure.



# Kube-proxy



# Installing kube-proxy

```
▶ wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kube-proxy
```

kube-proxy.service

```
ExecStart=/usr/local/bin/kube-proxy \\  
  --config=/var/lib/kube-proxy/kube-proxy-config.yaml  
Restart=on-failure  
RestartSec=5
```

## Three ways to see options

① Check in the `/etc/kubernetes/manifests`  
for YAML - for daemons

② Check in service folder.

`/etc/systemd/system/ - ....`

③ Check process

`ps -aux | grep <name>`